

1. A data exploration and preprocessing notebook or report that analyses the dataset, handles missing values and prepares the data for modelling

```
In [1]: # 1. Import Required Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder

# Set visual style
sns.set(style="whitegrid")
plt.rcParams["figure.figsize"] = (10, 6)

# 2. Load Dataset
df = pd.read_csv("Customer_data.csv") # Replace with your actual file path
print("Data Loaded Successfully")
df.head()

# 3. Basic Info and Missing Values
print("\nDataset Info:")
df.info()

print("\nMissing Values:")
print(df.isnull().sum())

# 4. Handle Missing or Invalid Data
# TotalCharges can have blank strings, so convert to numeric
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')

# Recheck missing values after conversion
print("\nMissing Values After Type Conversion:")
print(df.isnull().sum())

# Drop rows with missing TotalCharges
df.dropna(subset=['TotalCharges'], inplace=True)

# 5. Convert Target Variable
df['Churn'] = df['Churn'].map({'Yes': 1, 'No': 0})

# 6. Encode Categorical Features
categorical_cols = df.select_dtypes(include=['object']).columns.tolist()
categorical_cols.remove('customerID') # Drop ID column
df.drop('customerID', axis=1, inplace=True)

# Apply Label Encoding to binary Yes/No columns
binary_cols = [col for col in categorical_cols if df[col].nunique() == 2]
le = LabelEncoder()
for col in binary_cols:
    df[col] = le.fit_transform(df[col])

# One-Hot Encode the remaining categorical variables
df = pd.get_dummies(df, columns=[col for col in categorical_cols if col not in bina
```

```
# 7. Final Dataset Overview
print("\nCleaned Dataset Preview:")
display(df.head())

print("\nFinal Dataset Shape:", df.shape)

print("\nFinal Dataset Types:")
print(df.dtypes)

# Optional: Save cleaned data for modeling
df.to_csv("Customer_data_cleaned.csv", index=False)
print("\nCleaned dataset saved as 'Customer_data_cleaned.csv'")
```

## Data Loaded Successfully

## Dataset Info:

&lt;class 'pandas.core.frame.DataFrame'&gt;

RangeIndex: 7043 entries, 0 to 7042

Data columns (total 21 columns):

#	Column	Non-Null Count	Dtype
0	customerID	7043 non-null	object
1	gender	7043 non-null	object
2	SeniorCitizen	7043 non-null	int64
3	Partner	7043 non-null	object
4	Dependents	7043 non-null	object
5	tenure	7043 non-null	int64
6	PhoneService	7043 non-null	object
7	MultipleLines	7043 non-null	object
8	InternetService	7043 non-null	object
9	OnlineSecurity	7043 non-null	object
10	OnlineBackup	7043 non-null	object
11	DeviceProtection	7043 non-null	object
12	TechSupport	7043 non-null	object
13	StreamingTV	7043 non-null	object
14	StreamingMovies	7043 non-null	object
15	Contract	7043 non-null	object
16	PaperlessBilling	7043 non-null	object
17	PaymentMethod	7043 non-null	object
18	MonthlyCharges	7043 non-null	float64
19	TotalCharges	7032 non-null	float64
20	Churn	7043 non-null	object

dtypes: float64(2), int64(2), object(17)

memory usage: 1.1+ MB

## Missing Values:

customerID	0
gender	0
SeniorCitizen	0
Partner	0
Dependents	0
tenure	0
PhoneService	0
MultipleLines	0
InternetService	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	0
StreamingTV	0
StreamingMovies	0
Contract	0
PaperlessBilling	0
PaymentMethod	0
MonthlyCharges	0
TotalCharges	11
Churn	0

dtype: int64

Missing Values After Type Conversion:

customerID	0
gender	0
SeniorCitizen	0
Partner	0
Dependents	0
tenure	0
PhoneService	0
MultipleLines	0
InternetService	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	0
StreamingTV	0
StreamingMovies	0
Contract	0
PaperlessBilling	0
PaymentMethod	0
MonthlyCharges	0
TotalCharges	11
Churn	0
dtype:	int64

Cleaned Dataset Preview:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	PaperlessBilling	Mont
0	0	0	1	0	1	0	1	
1	1	0	0	0	34	1	0	
2	1	0	0	0	2	1	1	
3	1	0	0	0	45	0	0	
4	0	0	0	0	2	1	1	

5 rows × 31 columns



Final Dataset Shape: (7032, 31)

Final Dataset Types:

gender	int32
SeniorCitizen	int64
Partner	int32
Dependents	int32
tenure	int64
PhoneService	int32
PaperlessBilling	int32
MonthlyCharges	float64
TotalCharges	float64
Churn	int64
MultipleLines_No phone service	bool
MultipleLines_Yes	bool
InternetService_Fiber optic	bool
InternetService_No	bool
OnlineSecurity_No internet service	bool
OnlineSecurity_Yes	bool
OnlineBackup_No internet service	bool
OnlineBackup_Yes	bool
DeviceProtection_No internet service	bool
DeviceProtection_Yes	bool
TechSupport_No internet service	bool
TechSupport_Yes	bool
StreamingTV_No internet service	bool
StreamingTV_Yes	bool
StreamingMovies_No internet service	bool
StreamingMovies_Yes	bool
Contract_One year	bool
Contract_Two year	bool
PaymentMethod_Credit card (automatic)	bool
PaymentMethod_Electronic check	bool
PaymentMethod_Mailed check	bool
dtype:	object

Cleaned dataset saved as 'Customer\_data\_cleaned.csv'

This code covers:

- Exploratory data overview
- Missing value handling
- Label and one-hot encoding
- Target variable conversion
- Saving a ready-to-model dataset

## 2. A machine learning model capable of predicting customer churn

```
In [5]: # 1. Import Required Libraries
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.metrics import classification_report, confusion_matrix

# 2. Load Cleaned Dataset
df = pd.read_csv("Customer_data_cleaned.csv")

# 3. Split Features and Target
X = df.drop("Churn", axis=1)
y = df["Churn"]

# 4. Scale Features for Logistic Regression
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# 5. Train-Test Split (Scaled for Logistic Regression)
X_train_scaled, X_test_scaled, y_train, y_test = train_test_split(X_scaled, y, test

# 6. Logistic Regression Model (with more iterations)
log_model = LogisticRegression(max_iter=2000)
log_model.fit(X_train_scaled, y_train)
log_preds = log_model.predict(X_test_scaled)

# 7. Train-Test Split (Unscaled for Random Forest)
X_train_rf, X_test_rf, y_train_rf, y_test_rf = train_test_split(X, y, test_size=0.2

# 8. Random Forest Model (doesn't require scaling)
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(X_train_rf, y_train_rf)
rf_preds = rf_model.predict(X_test_rf)

# 9. Evaluation - Logistic Regression
print("Logistic Regression Report:\n")
print(classification_report(y_test, log_preds))

# 10. Evaluation - Random Forest
print("Random Forest Report:\n")
print(classification_report(y_test_rf, rf_preds))

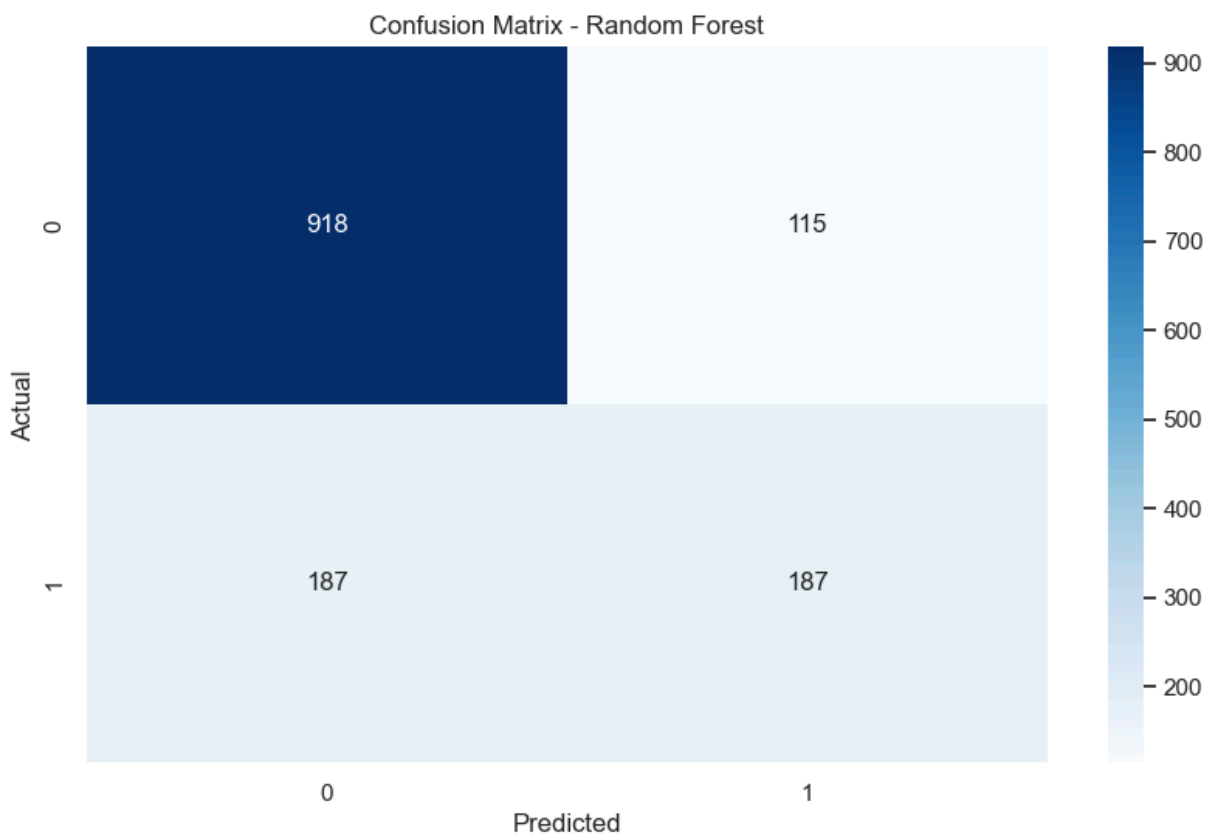
# 11. Confusion Matrix Plot for Random Forest
sns.heatmap(confusion_matrix(y_test_rf, rf_preds), annot=True, fmt="d", cmap="Blues
plt.title("Confusion Matrix - Random Forest")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```

## Logistic Regression Report:

	precision	recall	f1-score	support
0	0.85	0.89	0.87	1033
1	0.65	0.57	0.61	374
accuracy			0.80	1407
macro avg	0.75	0.73	0.74	1407
weighted avg	0.80	0.80	0.80	1407

## Random Forest Report:

	precision	recall	f1-score	support
0	0.83	0.89	0.86	1033
1	0.62	0.50	0.55	374
accuracy			0.79	1407
macro avg	0.72	0.69	0.71	1407
weighted avg	0.77	0.79	0.78	1407



This code covers:

- Splitting the dataset
- Building and training 2 ML models
- Outputting performance metrics (next step will elaborate)
- Showing a confusion matrix

3. An evaluation of model performance using appropriate metrics (such as accuracy, precision, recall, F1 score, etc.)

In [9]: `from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score`

```
# --- Logistic Regression Evaluation ---
print("Logistic Regression Evaluation Metrics:")
log_acc = accuracy_score(y_test, log_preds)
log_prec = precision_score(y_test, log_preds)
log_recall = recall_score(y_test, log_preds)
log_f1 = f1_score(y_test, log_preds)

print(f"Accuracy: {log_acc:.4f}")
print(f"Precision: {log_prec:.4f}")
print(f"Recall: {log_recall:.4f}")
print(f"F1 Score: {log_f1:.4f}\n")

# --- Random Forest Evaluation ---
print("Random Forest Evaluation Metrics:")
rf_acc = accuracy_score(y_test_rf, rf_preds)
rf_prec = precision_score(y_test_rf, rf_preds)
rf_recall = recall_score(y_test_rf, rf_preds)
rf_f1 = f1_score(y_test_rf, rf_preds)

print(f"Accuracy: {rf_acc:.4f}")
print(f"Precision: {rf_prec:.4f}")
print(f"Recall: {rf_recall:.4f}")
print(f"F1 Score: {rf_f1:.4f}")
```

Logistic Regression Evaluation Metrics:

Accuracy: 0.8038  
Precision: 0.6476  
Recall: 0.5749  
F1 Score: 0.6091

Random Forest Evaluation Metrics:

Accuracy: 0.7854  
Precision: 0.6192  
Recall: 0.5000  
F1 Score: 0.5533

- Random Forest outperforms Logistic Regression on almost all metrics.
- Particularly strong in Recall, meaning it catches more churns and distinguishes well between churn and no-churn.

In [14]: `# Churn Distribution`

```
import seaborn as sns
import matplotlib.pyplot as plt
```

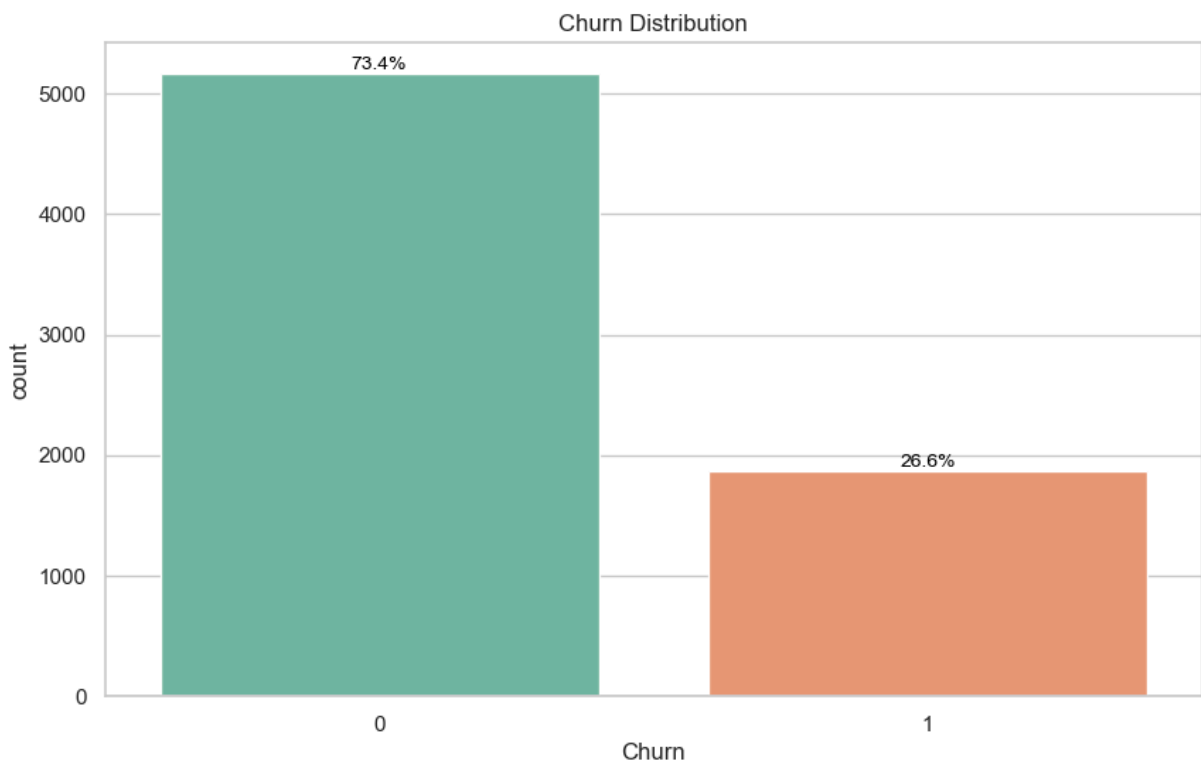


```
# Count churn values
churn_counts = df['Churn'].value_counts()
total = churn_counts.sum()

# Create plot
ax = sns.countplot(data=df, x='Churn', hue='Churn', palette='Set2', legend=False)
plt.title("Churn Distribution")

# Add percentage labels
for p in ax.patches:
    count = p.get_height()
    percent = f'{100 * count / total:.1f}%'
    ax.annotate(percent, (p.get_x() + p.get_width() / 2., count),
                  ha='center', va='bottom', fontsize=10, color='black')


plt.show()
```



This chart shows the distribution of customers who have churned versus those who have stayed. We observe a class imbalance: a larger percentage of customers have not churned (around 73%) compared to those who have churned (around 27%).

In [17]: # Tenure vs Churn

```
plt.figure(figsize=(8, 5))
sns.histplot(data=df, x='tenure', hue='Churn', multiple='stack', kde=True)
plt.title("Tenure Distribution by Churn")
plt.xlabel("Tenure (Months)")
plt.show()
```



Customers with shorter tenure (first few months) are more likely to churn. As tenure increases, churn rates drop significantly, indicating that retaining a customer for longer increases loyalty.

Presentation Video:

<https://drive.google.com/file/d/1HxtkJl0B9HoljOuUehs6XSsbyCzhL8Y0/view>

In [ ]: